Sindhu Marpureddy

TRUST & SAFETY | AI SAFETY & EVALUATION | FRAUD & RISK INVESTIGATIONS | QUALITY & POLICY

marpureddysindhu@gmail.com | +91 99630 51765 | Pune, India (open to relocation) | LinkedIn | GitHub | Portfolio

SUMMARY

Trust & Safety and AI evaluation professional with 13+ years of experience, including 10+ in Trust & Safety, at GitHub (Microsoft), Meta's programs and CrowdStrike. Investigates how platforms and AI systems get misused, writes the policies and playbooks that stop it, and builds the quality systems and tools that keep decisions consistent. Known for turning ambiguity into systems people can trust, and for building AI-assisted tools from scratch. B.Tech in Computer Science; fluent in English, Hindi, Telugu and Odia.

SELECTED IMPACT

- **Investigations at scale.** Resolved 500+ high-severity escalations a year at GitHub across fraud, account takeover, phishing, crypto fraud, CSAM, NCII and coordinated abuse; cut incident resolution time by 30%.
- **Evaluation that moved the numbers.** Designed evaluation frameworks on Meta's programs that improved review accuracy by 25% and cut high-risk turnaround by 20%.
- **AI evaluation.** Built a golden dataset and evaluated a production AI agent at CrowdStrike for accuracy, hallucinations, compliance and failure modes.
- **Policy for emerging harms.** Wrote policy and playbooks on synthetic media, email spoofing used for impersonation, and account takeover at GitHub.
- **Quality at scale.** Owned quality for a 108-member multilingual team; lifted CSAT by 5% and delivered 98% satisfaction with under 2% errors.

EXPERIENCE

Legal Analyst - AI Governance & Legal Operations — **CrowdStrike** | **June 2026 to October 2026**

- Evaluated AI agent outputs for accuracy and compliance in Salesforce, ensuring audit-ready documentation.
- Curated the RAG knowledge base to maintain the legal content library's accuracy and compliance.
- Developed **Sage**, an AI-assisted legal policy librarian using Google Apps Script, enhancing policy governance accuracy.
- Managed contract execution tracking in Ironclad, ensuring an accurate contract repository for stakeholders.

Senior Intelligence Consultant (Freelance) — **ALICE, formerly ActiveFence** | **March 2026 to April 2026**

- Delivered early-warning intelligence for a Rapid Alert Team covering India through OSINT, threat hunting and narrative forensics.

Trust and Safety Specialist II — **GitHub (Microsoft)** | **April 2021 to January 2026**

- Investigated account takeover, phishing, crypto fraud, scams, malware, CSAM, NCII, TVEC and coordinated abuse in Zendesk, resolving 500+ high-severity escalations a year across US and weekend shifts.
- Authored account-takeover SOPs and a remediation playbook, plus policy on synthetic media and email spoofing used for impersonation and copyright abuse.
- Reduced incident resolution time by 30% with Engineering, Legal and Security; turned abuse patterns into 50+ improvement plans with Threat Intelligence.
- Led calibration and edge-case adjudication as escalation SME for novel abuse patterns.

Process Developer (Trust and Safety Quality Analyst) — **Facebook via Spectraforce & Genpact** | **May 2015 to April 2021**

- Improved review accuracy by 25% and cut high-risk turnaround by 20% by designing evaluation frameworks measuring accuracy, precision, false-positive rate and TAT.
- Owned quality for a 108-member multilingual team and market quality for South Asia and North America, presenting to client leadership in weekly and monthly business reviews.
- Ran weekly calibration and error-review sessions across Operations, Policy and Quality; lifted bottom-quartile performers through the Revive program.
- Lifted CSAT by 5% by redesigning a moderation process; made automation 10% faster with a market-context knowledge repository.

Process Associate, Geo Operations — **Google Maps via Randstad / Minvesta** | **November 2012 to April 2015**

- Quality-controlled 2D/3D mapping and Street View imagery, including privacy blurring; created training material and an operator performance grid.

PROJECTS (BUILT INDEPENDENTLY)

Argus - real-time risk dashboard across 15 harm categories | **FatiGat** - reviewer fatigue and quality signals | **OctoMask** - privacy tool for developer profiles | **ReflectoCat** - tone and behavioral signals in GitHub activity

SKILLS

Trust & Safety: investigations, policy and playbooks, escalations, harms assessment, risk taxonomies | **AI:** AI agent evaluation, golden datasets, RAG knowledge bases, AI-assisted building | **Quality:** audits, calibration, evaluation frameworks, root-cause analysis | **Intelligence:** OSINT, threat hunting | **Regulation:** DSA, EU AI Act, GDPR | **Tools:** Zendesk, Salesforce, Ironclad, Google Apps Script, GitHub Copilot

EDUCATION

B.Tech, Computer Science Engineering, BPUT | CGPA 7.3 | **Lean Six Sigma Yellow Belt**